

DAG-GFlowNet

1 Bayesian Structure Learning

Definition 1. *Bayesian Network* (G, θ_G) is a probabilistic model over d random variables $\mathbf{X} = \{X_1, \dots, X_d\}$, whose joint distribution factorizes according to DAG G :

$$p(\mathbf{X}|G, \theta_G) = \prod_{i=1}^d p(X_i | \text{Pa}_G(X_i), \theta_i), \quad (1)$$

where $\text{Pa}_G(X_i)$ is the set of parents of the node X_i , θ_i parameters that define the probability distribution.

We are given a dataset $\mathcal{D} = (\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(N)})$ of N observations of random variable $\mathbf{X}$. The problem is to characterize the posterior distribution $P(G|D)$ over Bayesian Network structures that model these observations. Using Bayes formula, we can derive:

$$p(G|\mathcal{D}) = \frac{p(G, D)}{p(D)} = \frac{\int_{\theta_G} p(G, D, \theta_G) d\theta_G}{p(D)} = \frac{\int_{\theta_G} p(D|G, \theta_G) p(\theta_G|G) p(G) d\theta_G}{p(D)}. \quad (2)$$

Since we do not have access to $p(D)$, for any DAG G , we define its reward as the joint probability density function that equals to $p(G|D)$ up to constant $P(D)$:

$$R(G) = p(G, D) = \int_{\theta_G} p(D|G, \theta_G) p(\theta_G|G) p(G) d\theta_G. \quad (3)$$

Then, making assumptions of independence for parameter priors and of modularity for both parameter and DAG priors, we can derive closed form expression for the reward $R(G)$ as a product of local scores that only depend on individual variables and their parents in G :

$$R(G) = \prod_{i=1}^d \text{LocalScore}(X_i | \text{Pa}_G(X_i)). \quad (4)$$

Thus, we can calculate the reward for every DAG G . However, we still can not calculate the normalizing constant for the posterior since the number of possible DAGs is very large, making it computationally impossible. What we can do is to try to sample from the posterior $P(G|D)$.

To evaluate the quality of sampling method we will use *MAD*, i.e., *the maximum absolute deviation*:

$$\text{MAD} = \max_{i,j} |p_{ij} - \hat{p}_{ij}|, \quad (5)$$

where $\hat{p}_{ij}$ is the relative frequency of arc ij in the sampled DAGs and p_{ij} is the respective exact probability we compute by an exponential-time algorithm.

2 Experimental setup

To measure the performance of models on real datasets with less than 20 vertices, we calculate pearson and spearman correlation coefficients between true posterior probabilities $p_{ij} = P\{(i \rightarrow j) \in E(G)\}$, $G \sim P(G|D)$, computed via exponential-time algorithm (8), and empirical probabilities $\hat{p}_{ij}$, estimated using our sampler-model. Experimental results of vanilla DAG-GFlowNet on Asia dataset are presented on Figure 1.

Empirical probabilities for Asia dataset were estimated on 100K samples.




Figure 1: a) vanilla DAG-GFlowNet performance on Asia dataset; b) vanilla DAG-GFlowNet performance on Sachs dataset; c) DAG-GFlowNet with replay buffer performance on Sachs dataset; d) DAG-GFlowNet with replay buffer performance on Zoo dataset

References

- [1] X. Xie, Z. Geng, and Q. Zhao. Decomposition of structural learning about directed acyclic graphs. *Artificial Intelligence*, 170(4–5):422–439, 2006.
- [2] M. Zhu and S. Liu. A decomposition algorithm for learning Bayesian networks based on scoring function. *Journal of Applied Mathematics*, 2012:974063, 2012.
- [3] J. Dai, J. Ren, and W. Du. Decomposition-based Bayesian network structure learning algorithm using local topology information. *Knowledge-Based Systems*, 195:105602, 2020.

- [4] M. Wu, Y. Bao, R. Barzilay, and T. Jaakkola. Sample, estimate, aggregate: A recipe for causal discovery foundation models. *Transactions on Machine Learning Research*, 2025.
- [5] H. Sun, P. Tian, Z. Zhou, J. Zhang, P. Li, and A. L. Liu. Efficient causal structure learning via modular subgraph integration. *arXiv preprint arXiv:2601.21014*, 2026.
- [6] D. Nikzad, A. Zhilkin, J. Harviainen, J. Kuipers, G. Moffa, and M. Koivisto. Scaling up Bayesian DAG sampling. *arXiv preprint arXiv:2510.25254*, 2025.
- [7] J. Roy, P.-L. Bacon, C. Pal, and E. Bengio. Goal-conditioned GFlowNets for controllable multi-objective molecular design. In *ICML Workshop on Challenges in Deployable Generative AI*, 2023.
- [8] J. Pensar, T. Talvitie, A. Hyttinen, and M. Koivisto. A Bayesian approach for estimating causal effects from observational data. In *Proceedings of the AAAI conference on artificial intelligence* (Vol. 34, No. 04, pp. 5395-5402), 2020.